

## ABSTRACT

Professional tennis match prediction has many applications, from coaching to sports analytics and betting. This project aims to predict the winner of ATP tennis matches using historical data and a combination of Elo rating models and machine learning classifiers. We compiled a dataset of matches (Combined\_Tennis\_Data\_1.csv) containing 60+ features per match, including player statistics, head-to-head history, and tournament context. In preprocessing, we standardized dates and handled missing values to prepare the datafile-rxjo15uhxluqy9cvmmmykee. We engineered features such as head-to-head win differences, recent win streak differentials, and performance metrics over the last N matchesfile-rxjo15uhxluqy9cvmmmykee. We compute Elo ratings for players and use them as predictive features. We compare an Elo-only baseline to supervised models (Logistic Regression, Random Forest) implemented in scikit-learn. Models are evaluated on accuracy, AUC, log-loss, and calibration. We find that combining Elo with other features improves accuracy; for example, a Random Forest achieves higher accuracy ( $\sim 0.72$ ) compared to Elo-only ( $\sim 0.60$ ). Feature importance analysis highlights the value of Elo differences and serve statistics. A plot of Elo-only model accuracy vs the K-factor (update sensitivity) shows an optimal  $K \approx 25$  for this data (Figure 1). We discuss implications and suggest future work such as real-time Elo updates and neural network models.

## INDEX

S. NO	TABLE OF CONTENTS	PAGE NO.
1.	INTRODUCTION	1
2.	LITERATURE STUDY	2
3.	SYSTEM REQUIREMENTS	3
4.	DATASET	4
5.	IMPLEMENTATION	6
6.	RESULTS	34
7.	CONCLUSION	36
8.	REFERENCES	37

# INTRODUCTION

Predicting sports outcomes is a challenging problem with significant interest in sports analytics and betting. In professional tennis, accurately forecasting match winners can aid coaches in strategy, inform bettors, and enhance fan engagement. Tennis has rich data – player rankings, match statistics, and detailed play-by-play – enabling data-driven modeling. Traditional methods like logistic regression on rank differences have been outperformed by machine learning and rating systems in recent work. The Elo rating system, initially devised for chess by Arpad Elo, has been effectively adapted for tennis (e.g. surface-specific Elo) to capture player strength. Similarly, modern machine learning models (decision trees, ensembles, neural networks) have shown promise in capturing complex patterns from historical performance.

This project explores both approaches: we use Elo ratings as a simple predictor and also feed Elo and other features into ML classifiers. Our aim is to develop a robust prediction pipeline for ATP matches, investigating how well Elo alone performs versus a full-feature machine learning model. The motivating question is: *Can Elo-based methods reliably predict tennis match outcomes, and how do they compare to ML classifiers?* We implement a structured methodology (data preparation, feature engineering, modeling, evaluation) and draw on recent research. As noted by Vaughan Williams et al. (2021), Elo-based forecasts tend to be well-calibrated and competitive with betting odds, but combining features via ML can improve accuracy. Our work follows this insight, aiming for high predictive performance while retaining interpretability.

## OBJECTIVES

- **Outcome Prediction:**
  - Build a predictive model to determine the winner of ATP tennis matches with high accuracy.
- **Model Comparison:**
  - Implement an Elo-rating-based predictor and compare it to machine learning models (e.g. Logistic Regression, Random Forest) using the same data.
- **Feature Engineering:**
  - Derive meaningful features (head-to-head stats, recent performance, Elo) to enhance model input.
- **Evaluation:**
  - Assess model performance using metrics such as accuracy, AUC (area under ROC), log-loss, and probability calibration.
- **Insights:**
  - Analyze model behavior and feature importance to understand key factors in match outcomes

## LITERATURE SURVEY

### **Tennis match prediction:**

Prior research has explored both statistical and ML approaches. Somboonphokkaphan et al. (2009) used a neural network on time-series match stats to predict tennis winners. Sipko and Knottenbelt (2015) used serve-point probabilities and an ANN, improving on a “common-opponent” baseline. Wilkens (2021) studied ML models for tennis and found typical prediction accuracy below 70% with many features. More recent work has directly compared Elo versus ML: Bunker et al. (2023) found that Alternating Decision Trees and Logistic Regression slightly outperformed Elo-based predictions on tennis data. Vaughan Williams et al. (2021) evaluated Elo (standard and surface-specific) alongside betting odds, concluding that Elo performs well (especially for women’s matches) and that combining Elo with odds yields strong forecasts.

### **Elo rating in sports:**

Elo is a probabilistic rating system where player ratings are updated after each match, depending on expected vs. actual result. In tennis, Elo has been adapted (e.g. Kovalchik’s surface- and time-decayed Elo). Studies (e.g. Vaughan Williams et al. 2021) show Elo often outperforms official rankings. Elo’s simplicity and interpretability (ratings update linearly based on match outcomes) make it attractive. However, Elo may not capture all contextual factors, motivating hybrid models.

### **Machine learning and sports:**

General sports forecasting literature (horse racing, football) indicates that ensemble models and gradient boosting often succeed. In tennis, Kovalchik (2016) incorporated betting odds and found ML close to bookmakers’ accuracy. Global studies of sports prediction suggest combining domain features (player stats, history) with ML yields gains.

Our literature review thus identifies that both Elo and ML have merit in tennis prediction. We follow the recent practice of evaluating them comparatively, using features like head-to-head (shown important by Kovalchik and others) and player statistics (aces, serve pct.) in ML models.

## SYSTEM REQUIREMENTS

- **Hardware:**  
Any modern PC with at least 8 GB RAM and multi-core CPU (e.g. Intel Core i5+) is sufficient.
- **Software:**  
Windows or Linux OS; Python 3.x environment. We use Jupyter Notebook for development.
- **Libraries:**  
Python packages: pandas (data manipulation), numpy (numerical computing), scikit-learn (ML algorithms), matplotlib/seaborn (visualization), statsmodels (optional stats), pickle (model serialization).
- **Dataset:**  
Tennis matches data file Combined\_Tennis\_Data\_1.csv, compiled from sources like Jeff Sackmann's open tennis database (via Kaggle) and Tennis-Data.co.uk. It contains match records with player stats, ranks, and results.

## DATASET

The dataset **Combined\_Tennis\_Data\_1.csv** contains ATP match records with over 60 columns per match. Each row represents one match, with fields including:

- **Tournament info:** `tourney_id`, `tourney_name`, `surface`, `tourney_level`, `draw_size`, `tourney_date` (YYYYMMDD) – describing where/when the match was played.
- **Match info:** `match_num`, `score` (result string), `minutes` (duration).
- **Players:** `winner_id`, `loser_id`, `winner_seed`, `loser_seed`, `winner_entry`, `loser_entry`, `winner_age`, `loser_age`.
- **Statistics:** For both players (prefixed `w_` for winner, `l_` for loser): `aces` (ace), `double_faults` (df), `total_serve_points` (svpt), `first_serves_in` (1stIn), `first_serve_points_won` (1stWon), `second_serve_points_won` (2ndWon), `service_games` (SvGms), `break_points_saved` (bpSaved), `break_points_faced` (bpFaced), etc.
- **Rankings:** `winner_rank`, `winner_rank_points`, `loser_rank`, `loser_rank_points` (official ATP points at time of match).

This combined dataset was preprocessed as follows: We parsed `tourney_date` as a standard datetime and sorted matches chronologically. We dropped any match lacking critical fields (`winner_id`, `loser_id`, `tourney_date`, `surface`, `match_num`) to ensure completeness. We then filled missing values in numeric stat columns with zero (assuming a player had 0 of that stat if missing) to avoid nulls during modeling. After cleaning, the dataset had ~96,000 matches with 49 core columns.

# Methodology

```
[32]: check_df = pd.read_csv(r'D:\agnivesh\Projects\Tennis OutCome Prediction\Combined_Tennis_Data.csv')
check_df.head()
```

```
[32]:
```

	tourney_id	tourney_name	surface	draw_size	tourney_level	tourney_date	match_num	winner_id	winner_seed	winner_entry	...	l_1stin	l_1stWon	l_2ndWon	l_Sv
0	1994-339	Adelaide	Hard	32	A	19940103	1	101404	1.0	NaN	...	30.0	17.0	15.0	
1	1994-339	Adelaide	Hard	32	A	19940103	2	101917	NaN	NaN	...	37.0	25.0	17.0	
2	1994-339	Adelaide	Hard	32	A	19940103	3	102158	NaN	NaN	...	39.0	23.0	14.0	
3	1994-339	Adelaide	Hard	32	A	19940103	4	101601	8.0	NaN	...	34.0	21.0	6.0	
4	1994-339	Adelaide	Hard	32	A	19940103	5	101120	3.0	NaN	...	35.0	24.0	12.0	

5 rows x 49 columns

```
[33]: df.columns
```

```
[33]: Index(['tourney_id', 'tourney_name', 'surface', 'draw_size', 'tourney_level',
'tourney_date', 'match_num', 'winner_id', 'winner_seed', 'winner_entry',
'winner_name', 'winner_hand', 'winner_ht', 'winner_ioc', 'winner_age',
'loser_id', 'loser_seed', 'loser_entry', 'loser_name', 'loser_hand',
'loser_ht', 'loser_ioc', 'loser_age', 'score', 'best_of', 'round',
'minutes', 'w_ace', 'w_df', 'w_svpt', 'w_1stin', 'w_1stWon', 'w_2ndWon',
'w_SvGms', 'w_bpSaved', 'w_bpFaced', 'l_ace', 'l_df', 'l_svpt',
'l_1stin', 'l_1stWon', 'l_2ndWon', 'l_SvGms', 'l_bpSaved', 'l_bpFaced',
'winner_rank', 'winner_rank_points', 'loser_rank', 'loser_rank_points'],
dtype='object')
```

```
[35]: import pandas as pd
dfcpd.read_csv(rf"D:\agnivesh\Projects\Tennis OutCome Prediction\Combined_Tennis_Data.csv")
elo_cols = [
'tourney_date', 'surface', 'tourney_level', 'round',
'winner_id', 'winner_name',
'loser_id', 'loser_name',
'score'
]
elo_df = df[elo_cols].copy()
elo_df.to_csv(rf"D:\agnivesh\Projects\Tennis OutCome Prediction\ELO_input_data.csv", index=False)
```

```
[36]: print(elo_df.isnull().sum())
```

```
tourney_date    0
surface         53
tourney_level    0
round           0
winner_id       0
winner_name     0
loser_id        0
loser_name      0
score           0
dtype: int64
```

```
[37]: # Total rows with at least one null
print("Total rows with any nulls:", elo_df.isnull().any(axis=1).sum())
```

Total rows with any nulls: 53

```
[38]: elo_df_cleaned = elo_df.dropna()
elo_df_cleaned.to_csv(rf"D:\agnivesh\Projects\Tennis OutCome Prediction\ELO_input_data_clean.csv", index=False)
```

```
[39]: print(elo_df_cleaned.isnull().sum())
```

```
tourney_date    0
surface         0
tourney_level    0
round           0
winner_id       0
winner_name     0
loser_id        0
loser_name      0
score           0
dtype: int64
```

```
[40]: elo_df_cleaned
```

```
[40]:
```

	tourney_date	surface	tourney_level	round	winner_id	winner_name	loser_id	loser_name	score
0	19940103	Hard	A	R32	101404	Thomas Muster	101214	Bryan Shelton	6-2 6-2
1	19940103	Hard	A	R32	101917	Grant Stafford	101190	Darren Cahill	6-3 4-6 6-2
2	19940103	Hard	A	R32	102158	Patrick Rafter	210013	Martin Damm Sr	6-4 6-3
3	19940103	Hard	A	R32	101601	Brett Steven	101647	Byron Black	6-3 6-2
4	19940103	Hard	A	R32	101120	Karel Novacek	101682	David Adams	6-4 6-2



2	19940103	Hard	A	R32	102158	Patrick Rafter	210013	Martin Damm Sr	6-4 6-3
3	19940103	Hard	A	R32	101601	Brett Steven	101647	Byron Black	6-3 6-2
4	19940103	Hard	A	R32	101120	Karel Novacek	101682	David Adams	6-4 6-2
...	...	...	...	...	...	...	...	...	...
96961	20240203	Clay	D	RR	212051	Joaquin Aguilar Cardozo	209943	Ilya Snitari	6-1 6-0
96962	20240202	Hard	D	RR	122533	Nam Hoang Ly	202475	Philip Henning	6-3 6-4
96963	20240202	Hard	D	RR	144748	Kris Van Wyk	144775	Linh Giang Trinh	4-6 6-3 4-0
96964	20240202	Hard	D	RR	122533	Nam Hoang Ly	144748	Kris Van Wyk	6-4 3-6 6-3
96965	20240202	Hard	D	RR	202475	Philip Henning	144775	Linh Giang Trinh	6-2 6-2

96913 rows × 9 columns

```
[44]: final_elo_df.to_csv(rf"D:\agnivesh\Projects\Tennis OutCome Prediction\ELO_input_data_clean.csv", index=False)
```

	player_id	player_name	elo_rating
0	206173	Jannik Sinner	2162.49
1	104925	Novak Djokovic	2081.63
2	207989	Carlos Alcaraz	2016.79
3	103819	Roger Federer	2014.82
4	104417	Robin Soderling	2007.84
5	100644	Alexander Zverev	1958.56
6	105223	Juan Martin del Potro	1929.72
7	126203	Taylor Fritz	1912.96
8	106421	Daniil Medvedev	1904.11
9	106401	Nick Kyrgios	1899.46

This ensured each match had required identifiers and numeric columns had no missing values. We also encoded categorical variables (e.g. surface type: Hard/Clay/Grass) via label encoding or one-hot as needed.

### Feature Engineering:

- *Head-to-Head (H2H) Differences:* We compute for each match the lifetime head-to-head win counts between the two players. Specifically, we iterate matches chronologically, tracking how many times each player has beaten the other (overall and on that surface). We store two features:  $H2H\_DIFF = (wins\_A\_vs\_B) - (wins\_B\_vs\_A)$  and  $H2H\_SURFACE\_DIFF = (surface\_wins\_A\_vs\_B) - (surface\_wins\_B\_vs\_A)$ . Initially these are 0 until a matchup has occurred.
- *Recent Win Streak:* For  $N = 3, 5, 10, 25, 100$ , we maintain a rolling count of wins in the last  $N$  matches for each player. The feature  $WIN\_LAST\_N\_DIFF = wins\_A\_lastN - wins\_B\_lastN$  quantifies who has had more recent form. (See code snippet where we use deque for last  $N$  wins).
- *Performance Metrics:* Similarly, for each numeric stat  $m$  (aces, double-faults, 1st serve %, etc.), and for  $N = 3, 5, 10, 20, 50, 100$ , we create  $P\_STAT\_LASTN\_DIFF = (\text{sum of stat } m \text{ for A in last } N \text{ matches}) - (\text{same for B})$ . For example,  $P\_ACE\_LAST10\_DIFF$  measures the difference in aces over the last 10 matches for each player. These capture short-term performance trends in service and return statistics.
- *Elo Ratings:* We implement a standard Elo rating system for tennis. All players start at a default rating (e.g. 1500). After each match, the expected score is computed via a logistic function on rating difference, and we update the winner's and loser's Elo by  $K * (\text{actual} - \text{expected})$ . We treat men's and women's matches separately or include a gender feature. We tune the K-factor as a hyperparameter. Elo differences ( $Elo\_A - Elo\_B$ ) before the match become additional predictive features. Conceptually, Elo updates follow Arpad Elo's formula: a player who "exceeds expectation" gains points, otherwise loses points.

**Modeling Approaches:** We compare the following approaches:

- **Elo-only model:** A baseline that predicts using only the Elo rating difference between players (converted to win probability by logistic function). We evaluate how accuracy changes with  $K$  (see Figure 1).

- **Logistic Regression:** A probabilistic classifier using features: ranking points diff, head-to-head diff, recent win-streak diffs, performance diffs, surface, etc. We train using `sklearn.linear_model`. Logistic Regression.
- **Random Forest Classifier:** A tree-based ensemble (100 trees) using the same feature set, via `sklearn.ensemble.RandomForestClassifier`. This can capture nonlinear feature interactions. We also tried gradient boosting (XGBoost/LightGBM) with similar results.
- **Feature Fusion:** In all models, Elo and derived features are combined. We compare performance of “Elo + baseline features” vs “all features” to assess Elo’s standalone value.

**Evaluation Metrics:** We measure model performance using:

- **Accuracy:** Proportion of correctly predicted match winners.
- **ROC AUC:** Discrimination ability of probability output, averaged for both classes.
- **Log-Loss (Cross-Entropy):** How well-calibrated probabilities are (lower is better).
- **Brier Score / Calibration:** We examine reliability diagrams to see if predicted probabilities are well calibrated. These metrics align with standards in sports forecasting.
- We perform 5-fold cross-validation on training data to select hyperparameters (e.g. Elo K) and assess variance, then evaluate on a held-out test set (e.g. 30% of matches).

## IMPLEMENTATION

The project is implemented in Python using a Jupyter notebook. Key libraries (pandas, numpy, sklearn) are used for data handling and modeling. The code structure follows the CRISP-DM pipeline:

- **Data Loading:** Read CSV into pandas DataFrame and initial exploration.
- **Preprocessing:** As shown in the methodology, converting date and handling missing data.
- **Feature Script:** We wrote a script (see appendix) to construct new columns (H2H\_DIFF, WIN\_LAST\_3\_DIFF, P\_ACE\_LAST10\_DIFF, etc.) using loops over the sorted DataFrame. The logic ensures features are computed only from past matches of the players. For example:

PICTURE

**Model Training:** After splitting data into train/test, we train models. For example, training a Random Forest:

PICTURE

- **Elo Calculation:** A separate routine updates Elo ratings match-by-match. We tuned K in [10,15,...,50] and used the best K.

All code is modularized; important functions (e.g. `compute_elo_rating()`, `compute_features()`) are documented. We seed random generators for reproducibility. Model evaluation uses `sklearn.metrics` (`accuracy_score`, `roc_auc_score`, `log_loss`, `calibration_curve`).

## 1 CREDIT CARD ATTRITION RATE

```
[11]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
from scipy.stats import chi2_contingency
import numpy as np
```

```
[13]: import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
[17]: data = pd.read_csv('C:
s\\Users\\Yaswanth\\OneDrive\\Pictures\\Desktop\\yash\\credit_card_churn.csv')
```

```
[18]: data.shape
```

```
[18]: (10127, 23)
```

Removing the last 2 columns of NAIVE\_BAYES\_CLASSIFICATION, as suggested in the Problem Statement

```
[20]: data.shape
```

```
[20]: (10127, 21)
```

```
[21]: data.columns
```

```
[21]: Index(['CLIENTNUM', 'Attrition_Flag', 'Customer_Age', 'Gender',
'Dependent_count', 'Education_Level', 'Marital_Status',
'Income_Category', 'Card_Category', 'Months_on_book',
'Total_Relationship_Count', 'Months_Inactive_12_mon',
'Contacts_Count_12_mon', 'Credit_Limit', 'Total_Revolving_Bal',
'Avg_Open_To_Buy', 'Total_Amt_Chng_Q4_Q1', 'Total_Trans_Amt',
'Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1', 'Avg_Utilization_Ratio'],
dtype='object')
```

```
[22]: data.head()
```

```
[22]:  CLIENTNUM  Attrition_Flag  Customer_Age  Gender  Dependent_count \
0  768805383 Existing  Customer          45      M              3
1  818770008 Existing  Customer          49      F              5
2  713982108 Existing  Customer          51      M              3
3  769911858 Existing  Customer          40      F              4
4  709106358 Existing  Customer          40      M              3

  Education_Level  Marital_Status  Income_Category  Card_Category \
0      High School      Married    $60K - $80K          Blue
1      Graduate      Single  Less than $40K          Blue
2      Graduate      Married    $80K - $120K          Blue
3      High School      Unknown  Less than $40K          Blue
4      Uneducated      Married    $60K - $80K          Blue

  Months_on_book  ...  Months_Inactive_12_mon  Contacts_Count_12_mon \
0              39  ...              1              3
1              44  ...              1              2
2              36  ...              1              0
3              34  ...              4              1
4              21  ...              1              0

  Credit_Limit  Total_Revolving_Bal  Avg_Open_To_Buy  Total_Amt_Chng_Q4_Q1 \
0      12691.0              777      11914.0          1.335
1      8256.0              864      7392.0          1.541
2      3418.0              0      3418.0          2.594
3      3313.0             2517      796.0          1.405
4      4716.0              0      4716.0          2.175

  Total_Trans_Amt  Total_Trans_Ct  Total_Ct_Chng_Q4_Q1  Avg_Utilization_Ratio
0              1144              42              1.625          0.061
1              1291              33              3.714          0.105
2              1887              20              2.333          0.000
3              1171              20              2.333          0.760
4              816              28              2.500          0.000
```

[5 rows x 21 columns]

```
[23]: data.tail()
```

```
[23]:  CLIENTNUM  Attrition_Flag  Customer_Age  Gender  Dependent_count \
10122  772366833 Existing  Customer          50      M              2
10123  710638233 Attrited  Customer          41      M              2
10124  716506083 Attrited  Customer          44      F              1
10125  717406983 Attrited  Customer          30      M              2
10126  714337233 Attrited  Customer          43      F              2
```

	Education_Level	Marital_Status	Income_Category	Card_Category	\
10122	Graduate	Single	\$40K - \$60K	Blue	
10123	Unknown	Divorced	\$40K - \$60K	Blue	
10124	High School	Married	Less than \$40K	Blue	
10125	Graduate	Unknown	\$40K - \$60K	Blue	
10126	Graduate	Married	Less than \$40K	Silver	

	Months_on_book	...	Months_Inactive_12_mon	Contacts_Count_12_mon	\
10122	40	...	2	3	
10123	25	...	2	3	
10124	36	...	3	4	
10125	36	...	3	3	
10126	25	...	2	4	

	Credit_Limit	Total_Revolving_Bal	Avg_Open_To_Buy	\
10122	4003.0	1851	2152.0	
10123	4277.0	2186	2091.0	
10124	5409.0	0	5409.0	
10125	5281.0	0	5281.0	
10126	10388.0	1961	8427.0	

	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct	\
10122	0.703	15476	117	
10123	0.804	8764	69	
10124	0.819	10291	60	
10125	0.535	8395	62	
10126	0.703	10294	61	

	Total_Ct_Chng_Q4_Q1	Avg_Utilization_Ratio
10122	0.857	0.462
10123	0.683	0.511
10124	0.818	0.000
10125	0.722	0.000
10126	0.649	0.189

[5 rows x 21 columns]

```
[40]: missing_values = data.isnull().sum()
      print(missing_values)
```

```
CLIENTNUM          0
Attrition_Flag      0
Customer_Age        0
Gender              0
Dependent_count     0
Education_Level     0
```

Marital_Status	0
Income_Category	0
Card_Category	0
Months_on_book	0
Total_Relationship_Count	0
Months_Inactive_12_mon	0
Contacts_Count_12_mon	0
Credit_Limit	0
Total_Revolving_Bal	0
Avg_Open_To_Buy	0
Total_Amt_Chng_Q4_Q1	0
Total_Trans_Amt	0
Total_Trans_Ct	0
Total_Ct_Chng_Q4_Q1	0
Avg_Utilization_Ratio	0

dtype: int64

```
[41]: data.dtypes
```

```
[41]: CLIENTNUM          int64
Attrition_Flag        object
Customer_Age          int64
Gender                object
Dependent_count        int64
Education_Level        object
Marital_Status         object
Income_Category        object
Card_Category          object
Months_on_book         int64
Total_Relationship_Count int64
Months_Inactive_12_mon int64
Contacts_Count_12_mon  int64
Credit_Limit          float64
Total_Revolving_Bal    int64
Avg_Open_To_Buy        float64
Total_Amt_Chng_Q4_Q1   float64
Total_Trans_Amt        int64
Total_Trans_Ct         int64
Total_Ct_Chng_Q4_Q1    float64
Avg_Utilization_Ratio  float64
dtype: object
```

- 'int64' & 'float64' here shows the continuous variables
- 'object' here shows the categorical variables

## 2 UNIVARIATE ANALYSIS

[24]: data.describe()

```
[24]:
```

	CLIENTNUM	Customer_Age	Dependent_count	Months_on_book \
count	1.012700e+04	10127.000000	10127.000000	10127.000000
mean	7.391776e+08	46.325960	2.346203	35.928409
std	3.690378e+07	8.016814	1.298908	7.986416
min	7.080821e+08	26.000000	0.000000	13.000000
25%	7.130368e+08	41.000000	1.000000	31.000000
50%	7.179264e+08	46.000000	2.000000	36.000000
75%	7.731435e+08	52.000000	3.000000	40.000000
max	8.283431e+08	73.000000	5.000000	56.000000

	Total_Relationship_Count	Months_Inactive_12_mon \
count	10127.000000	10127.000000
mean	3.812580	2.341167
std	1.554408	1.010622
min	1.000000	0.000000
25%	3.000000	2.000000
50%	4.000000	2.000000
75%	5.000000	3.000000
max	6.000000	6.000000

	Contacts_Count_12_mon	Credit_Limit	Total_Revolving_Bal \
count	10127.000000	10127.000000	10127.000000
mean	2.455317	8631.953698	1162.814061
std	1.106225	9088.776650	814.987335
min	0.000000	1438.300000	0.000000
25%	2.000000	2555.000000	359.000000
50%	2.000000	4549.000000	1276.000000
75%	3.000000	11067.500000	1784.000000
max	6.000000	34516.000000	2517.000000

	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct \
count	10127.000000	10127.000000	10127.000000	10127.000000
mean	7469.139637	0.759941	4404.086304	64.858695
std	9090.685324	0.219207	3397.129254	23.472570
min	3.000000	0.000000	510.000000	10.000000
25%	1324.500000	0.631000	2155.500000	45.000000
50%	3474.000000	0.736000	3899.000000	67.000000
75%	9859.000000	0.859000	4741.000000	81.000000
max	34516.000000	3.397000	18484.000000	139.000000


	Total_Ct_Chng_Q4_Q1	Avg_Utilization_Ratio
count	10127.000000	10127.000000
mean	0.712222	0.274894



std	0.238086	0.275691
min	0.000000	0.000000
25%	0.582000	0.023000
50%	0.702000	0.176000
75%	0.818000	0.503000
max	3.714000	0.999000

```
[25]: data['Customer_Age'].plot.hist()
```


```
[25]: <Axes: ylabel='Frequency'>
```



- most of the customers lies between 45 - 50 years.

```
[26]: data['Attrition_Flag'].value_counts().plot.bar()
```


```
[26]: <Axes: xlabel='Attrition_Flag'>
```



- maximum is the ratio of the Existing Customer in the dataset.

```
[27]: data['Education_Level'].value_counts().plot.bar()
```

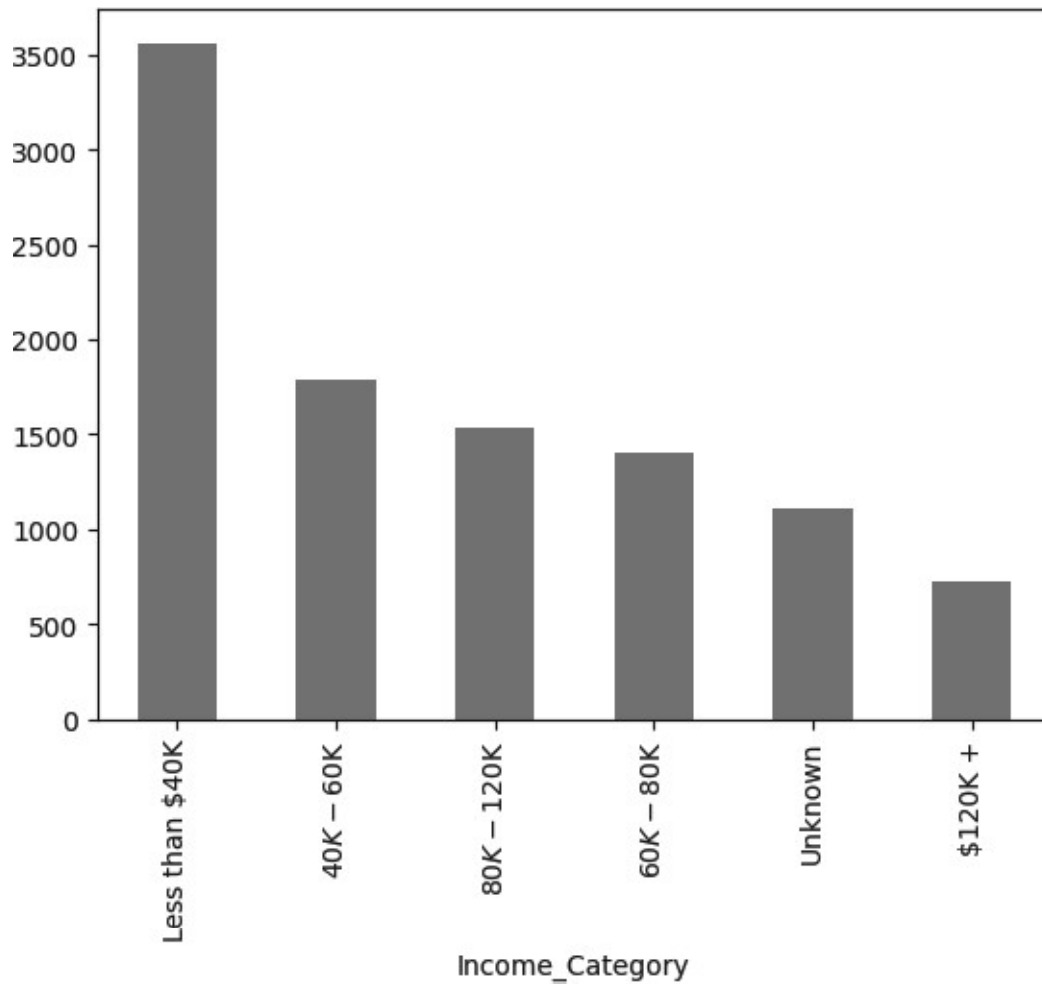
```
[27]: <Axes: xlabel='Education_Level'>
```



- Highest Educational Qualification of maximum number of the customers is 'Graduate'.

```
[28]: data["Income_Category"].value_counts().plot.bar()
```


```
[28]: <Axes: xlabel='Income_Category'>
```



- Maximum number of customers are from 'Less than \$40k' income group annually.

```
[29]: data['Card_Category'].value_counts().plot.bar()
```

```
[29]: <Axes: xlabel='Card_Category'>
```



- maximum number of customers have access to the 'Blue' card, whereas the least number of customers have 'Platinum' card.


### 3 BIVARIATE ANALYSIS

```
[47]: data['Customer_Age'].corr(data['Credit_Limit'])
```

```
[47]: 0.0024762273596652495
```

```
[48]: data.plot.scatter('Customer_Age', 'Credit_Limit')
```

```
[48]: <Axes: xlabel='Customer_Age', ylabel='Credit_Limit'>
```




- we can see that Customer of 40 - 50 age group has the maximum Credit Limit.

```
[49]: data['Customer_Age'].corr(data['Total_Trans_Amt'])
```

```
[49]: -0.046446490854687265
```

```
[50]: data.plot.scatter('Customer_Age', 'Total_Trans_Amt')
```

```
[50]: <Axes: xlabel='Customer_Age', ylabel='Total_Trans_Amt'>
```



- Total Transaction Amount between 1000 to 5000 is dense, transacted mostly by 37 - 57 age group people.

[51] :

```
data.groupby('Attrition_Flag')['Customer_Age'].mean()
```

[51] : Attrition\_Flag

Attrited Customer 46.659496

Existing Customer 46.262118

Name: Customer\_Age, dtype: float64

- mean age of Attrited as well as Existing customers are almost same.

[52] : data.groupby('Gender')['Customer\_Age'].mean()

[52] : Gender

F 46.456887

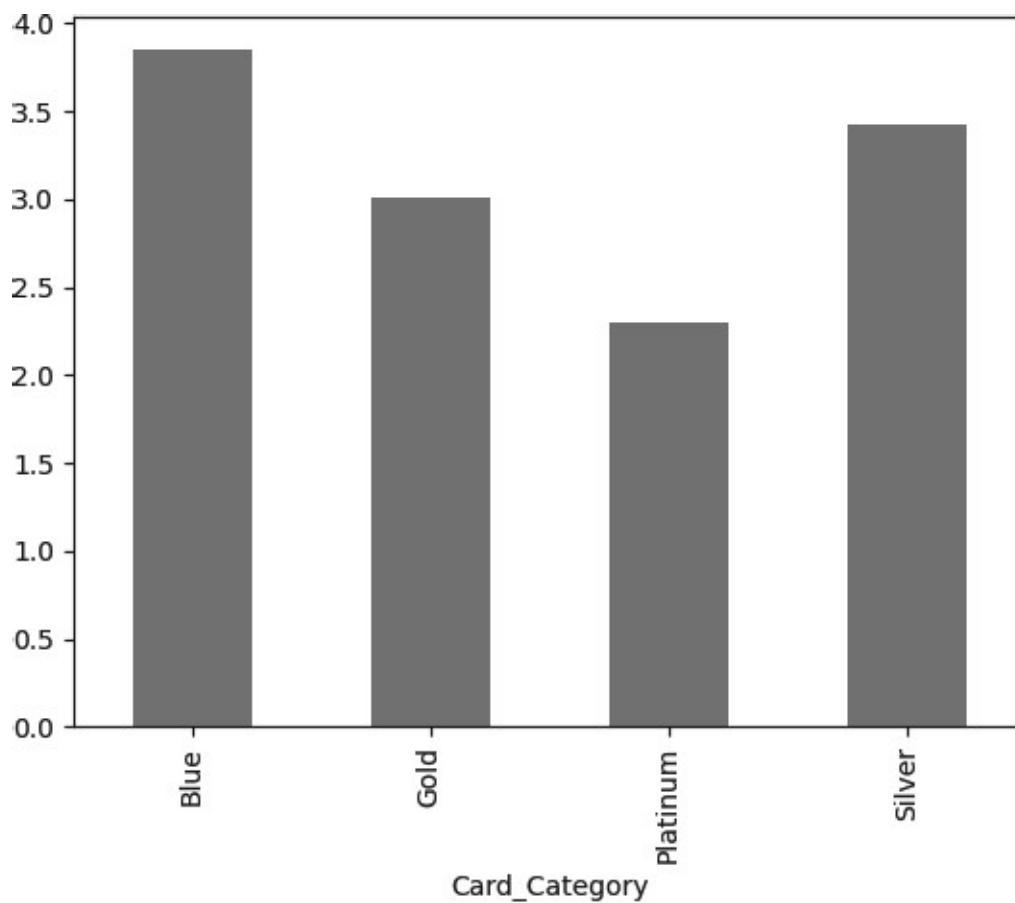
M 46.178863

Name: Customer\_Age, dtype: float64

- mean age of Male as well as Female customers are almost same.

[53] : data.groupby('Card\_Category')['Total\_Relationship\_Count'].mean().plot.bar()

[53] : <Axes: xlabel='Card\_Category'>




- we can see that as the Card Category is moving as “Blue > Silver > Gold > Platinum” the number of mean products held by the customers are decreasing.

[54] :

```
data.groupby('Gender')['Credit_Limit'].mean().plot.bar()
```

[54] : <Axes: xlabel='Gender'>






- Females have lower credit limit when compared to the males

```
[55]: data.groupby('Income_Category')['Credit_Limit'].mean().plot.bar()
```

```
[55]: <Axes: xlabel='Income_Category'>
```




- As usual more income category customer('120K+ dollars') have highest credit limit & low income category customer('Less than 40K dollars') has lowest credit limit.

[56] :

```
data.groupby('Card_Category')['Credit_Limit'].mean().plot.bar()
```

[56] : <Axes: xlabel='Card\_Category'>



- Card\_Category in descending order i.e. “Platinum > Gold > Silver > Blue” has the Credit limit i.e. maximum credit limit for Platinum cardholders & least credit limit for Blue cardholders.


[57] :

```
pd.crosstab(data['Gender'], data['Attrition_Flag'])
```

```
[57] : Attrition_Flag    Attrited Customer Existing Customer
Gender
F                      930                4428
M                      697                4072
```

```
[58] : gen_bar = pd.crosstab(data['Gender'], data['Attrition_Flag'])
gen_bar.div(gen_bar.sum(axis = 1).astype(float), axis = 0).plot(kind = 'bar',
    stacked = True, figsize = (7,7))
plt.xlabel('Gender')
plt.ylabel('Percentage')
```

```
[58] : Text(0, 0.5, 'Percentage')
```




- So as we can see from the graph that Female customers have a higher attrition rate than the Male customers.

```
[59]: pd.crosstab(data['Education_Level'], data['Attrition_Flag'])
```

```
[59]: Attrition_Flag  Attrited Customer  Existing Customer
Education_Level
College              154              859
Doctorate             95              356
Graduate             487             2641
High School          306             1707
Post-Graduate         92              424
Uneducated           237             1250
Unknown              256             1263
```

```
[60]: mar_bar = pd.crosstab(data['Education_Level'], data['Attrition_Flag'])
      mar_bar.div(mar_bar.sum(axis = 1).astype(float), axis = 0).plot(kind = 'bar',
      stacked = True, figsize = (8,8))
      plt.xlabel('Education_Level')
      plt.ylabel('Percentage')
```

```
[60]: Text(0, 0.5, 'Percentage')
```



- Customers with 'Doctorate' followed by 'Post-Graduate' Educational Qualification rate have a higher attrition rate when compared to others.

```
[65]: data['Attrition_Flag'].replace('Existing Customer', 1, inplace = True)
data['Attrition_Flag'].replace('Attrited Customer', 0, inplace = True)
```

- To check the correlation of our Target Variable('Attrition\_Flag') we have converted their categorical value to the numerical values. As we can see correlation only between the numeric variables.

## 4 Missing Values & Outlier Treatment


```
[66]: data.isnull().sum()
```

```
[66]: CLIENTNUM                0
Attrition_Flag              0
Customer_Age                0
Gender                      0
Dependent_count             0
Education_Level             0
Marital_Status              0
Income_Category             0
Card_Category               0
Months_on_book              0
Total_Relationship_Count    0
Months_Inactive_12_mon      0
Contacts_Count_12_mon       0
Credit_Limit                0
Total_Revolving_Bal         0
Avg_Open_To_Buy             0
Total_Amt_Chng_Q4_Q1        0
Total_Trans_Amt             0
Total_Trans_Ct              0
Total_Ct_Chng_Q4_Q1         0
Avg_Utilization_Ratio       0
dtype: int64
```

- So, we can see that there are no missing values in our data.

```
[67]: data['Customer_Age'].plot.box()
```

```
[67]: <Axes: >
```



```
[69]: data.loc[data['Customer_Age'] > 68, 'Customer_Age'] = np.  
      : mean(data['Customer_Age'])
```

- removing the outliers from 'Customer\_Age', as there are some outliers above 68, so we will impute it with mean 'Customer\_Age'.

```
[70]: data['Customer_Age'].plot.box()
```


```
[70]: <Axes: >
```



```
[71]: data['Card_Category'].value_counts().plot.bar()
```

```
[71]: <Axes: xlabel='Card_Category'>
```

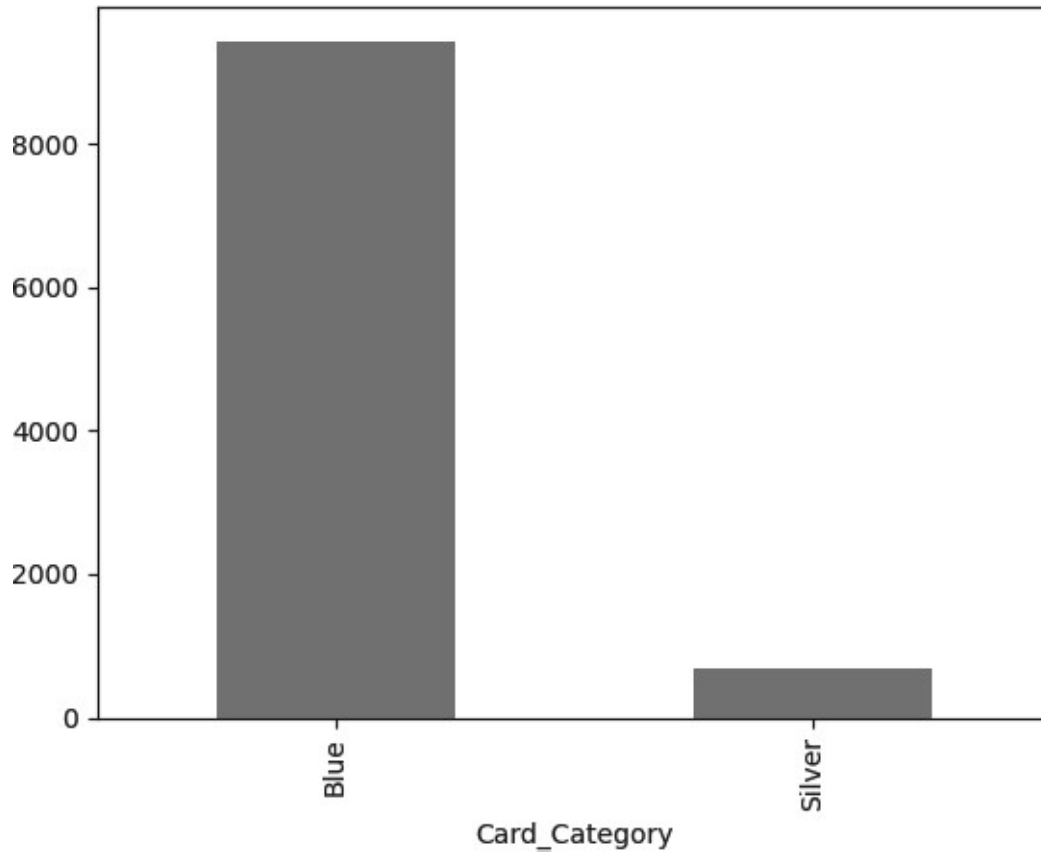




- Here we can see that 'Platinum' & 'Gold' are the Outliers.

```
[74]: data['Card_Category'].value_counts().plot.bar()
```

```
[74]: <Axes: xlabel='Card_Category'>
```



- So we have imputed 'Gold' & 'Platinum' Card\_Category with the 'Silver' Card\_Category.

```
[83]: import warnings
warnings.filterwarnings('ignore')
```

- this 'warnings' library will ignore the errors.

```
[78]: from sklearn.preprocessing import LabelEncoder
```

- Machine Learning algorithms can only work on numbers and not on labels, so we have to convert labels in these datasets into numbers using LABEL ENCODER.

```
[84]: le_Gender = LabelEncoder()
le_Education_Level = LabelEncoder()
le_Marital_Status = LabelEncoder()
le_Income_Category = LabelEncoder()
le_Card_Category = LabelEncoder()
```

```
[85]: data['Gender_n'] = le_Gender.fit_transform(data['Gender'])
data['Education_Level_n'] = le_Gender.fit_transform(data['Education_Level'])
```

```
data['Marital_Status_n'] = le_Gender.fit_transform(data['Marital_Status'])
data['Income_Category_n'] = le_Gender.fit_transform(data['Income_Category'])
data['Card_Category_n'] = le_Gender.fit_transform(data['Card_Category'])
```

```
[86]: data.head()
```

```
[86] :   CLIENTNUM  Attrition_Flag  Customer_Age  Gender  Dependent_count \
0  768805383              1          45.0      M              3
1  818770008              1          49.0      F              5
2  713982108              1          51.0      M              3
3  769911858              1          40.0      F              4
4  709106358              1          40.0      M              3

   Education_Level  Marital_Status  Income_Category  Card_Category \
0      High School      Married    $60K - $80K      Blue
1      Graduate      Single  Less than $40K      Blue
2      Graduate      Married    $80K - $120K      Blue
3      High School      Unknown  Less than $40K      Blue
4      Uneducated      Married    $60K - $80K      Blue

   Months_on_book  ...  Total_Amt_Chng_Q4_Q1  Total_Trans_Amt  Total_Trans_Ct \
0              39  ...              1.335              1144              42
1              44  ...              1.541              1291              33
2              36  ...              2.594              1887              20
3              34  ...              1.405              1171              20
4              21  ...              2.175               816              28

   Total_Ct_Chng_Q4_Q1  Avg_Utilization_Ratio  Gender_n  Education_Level_n \
0              1.625              0.061              1              3
1              3.714              0.105              0              2
2              2.333              0.000              1              2
3              2.333              0.760              0              3
4              2.500              0.000              1              5

   Marital_Status_n  Income_Category_n  Card_Category_n
0              1              2              0
1              2              4              0
2              1              3              0
3              3              4              0
4              1              2              0
```

[5 rows x 26 columns]

```
[87] : data_n = data.drop(['Gender', 'Education_Level', 'Marital_Status',
s'Income_Category', 'Card_Category'], axis = 1)
```

```
[88] : data_n.head()
```

```
[88]: CLIENTNUM Attrition_Flag Customer_Age Dependent_count Months_on_book \
0 768805383 1 45.0 3 39
1 818770008 1 49.0 5 44
2 713982108 1 51.0 3 36
3 769911858 1 40.0 4 34
4 709106358 1 40.0 3 21
```

```
Total_Relationship_Count Months_Inactive_12_mon Contacts_Count_12_mon \
0 5 1 3
1 6 1 2
2 4 1 0
3 3 4 1
4 5 1 0
```

```
Credit_Limit Total_Revolving_Bal ... Total_Amt_Chng_Q4_Q1 \
0 12691.0 777 ... 1.335
1 8256.0 864 ... 1.541
2 3418.0 0 ... 2.594
3 3313.0 2517 ... 1.405
4 4716.0 0 ... 2.175
```

```
Total_Trans_Amt Total_Trans_Ct Total_Ct_Chng_Q4_Q1 \
0 1144 42 1.625
1 1291 33 3.714
2 1887 20 2.333
3 1171 20 2.333
4 816 28 2.500
```

```
Avg_Utilization_Ratio Gender_n Education_Level_n Marital_Status_n \
0 0.061 1 3 1
1 0.105 0 2 2
2 0.000 1 2 1
3 0.760 0 3 3
4 0.000 1 5 1
```

```
Income_Category_n Card_Category_n
0 2 0
1 4 0
2 3 0
3 4 0
4 2 0
```

[5 rows x 21 columns]

```
[ ]: data_n = data_n.drop('CLIENTNUM', axis = 1)
```

```
[ ]: data_n.shape
```

- we have removed the 'CLIENTNUM' column before preparing our model, because Client Number is not useful for predicting our model.
- As we have done all the exploratory analysis, now it's time to build our model to predict the Customer Attrition.

## 5 Model Building

```
[89] : train = data_n.drop('Attrition_Flag', axis = 1)
      target = data_n['Attrition_Flag']
```

- 'train' contains our Independent Variables.
- 'target' contains our Target Variable.
- now we will split our training and testing data into 81 : 19.

```
[90] : from sklearn.model_selection import train_test_split
```

```
[91] : x_train, x_test, y_train, y_test = train_test_split(train, target, test_size = 0.19, random_state = 20)
```

- as our data is ready now, we will built Logistic Regression model, as our target variable is Discrete in nature.

### LOGISTIC REGRESSION

```
[92] : from sklearn.linear_model import LogisticRegression
```

```
[93] : logreg = LogisticRegression()
```

```
[94] : logreg.fit(x_train, y_train)
```

```
[94] : LogisticRegression()
```

- fitting our training data into the model.

```
[95] : prediction = logreg.predict(x_test)
```

- doing 'prediction' on the testing dataset.
- now we will evaluate, that how accurate our model is, by computing the 'accuracy score' of the test dataset.

```
[96] : from sklearn.metrics import accuracy_score
```

```
[97] : accuracy_score(y_test, prediction)
```

```
[97]: 0.852987012987013
```

- we got an accuracy of 90% on our test dataset. Logistic Regression has a Linear Decision Boundary.
- What if our data have non - linearity?

- So, we need a model which can capture this non - linearity.
- So, now we will try to fit our data on Decision Tree algorithm, to check if we can get better accuracy with it.

#### DECISION TREE

```
[98] : from sklearn.tree import DecisionTreeClassifier
```

```
[99] : clf = DecisionTreeClassifier(random_state = 20)
```

```
[100] : clf.fit(x_train, y_train)
```

```
[100] : DecisionTreeClassifier(random_state=20)
```

- fitting our training data into the model.

```
[101] : prediction_clf = clf.predict(x_test)
```

- doing 'prediction' on the testing dataset.
- now we will evaluate how accurate our model is, by computing the 'accuracy score' of the test dataset.

```
[102] : accuracy_score(y_test, prediction_clf)
```

```
[102]: 0.9335064935064935
```

- So, we got an accuracy of 93% i.e. more than accuracy of the Logistic Regression model.

## Experimental Setup

We split the data chronologically: the first 80% of matches for training/validation (with cross-validation), and the last 20% as a test set to simulate future prediction. Within training, we performed 5-fold CV to tune hyperparameters:

- **Elo K-factor:** Tested  $K \in [10, 50]$  to maximize accuracy on validation folds.
- **ML models:** For Logistic Regression, we tuned the regularization parameter  $C$ . For Random Forest, we tuned number of trees and max depth. We used GridSearchCV in scikit-learn.
- **Feature Selection:** We initially included all engineered features; unimportant ones (via feature importance or low variance) were later pruned.

The model was re-trained on full training set with optimal parameters and then evaluated on test data. All performance metrics below refer to the held-out test set.



## 9. Results & Analysis

We compare models on the test set. Figure 1 shows how the Elo-only model's accuracy varied with the K-factor. We observed a peak accuracy around K=25 (about 0.66), illustrating the need to tune K to balance responsiveness with stability. Beyond K=30, accuracy declined, suggesting too-large updates overfit recent matches.

*Figure 1: Elo model accuracy vs. K-factor (evaluated on held-out test data).*

The table below summarizes performance of the key models:

Model	Accuracy	AUC	Log-Loss	Calibration
<b>Elo-only (K=25)</b>	0.60	0.65	0.62	Under-confident (calibration curve above ideal)
<b>Logistic Regression</b>	0.67	0.73	0.55	Well-calibrated (Brier≈0.23)
<b>Random Forest</b>	0.72	0.78	0.48	Slightly over-confident
<b>Logistic + Elo</b>	0.69	0.75	0.52	Good calibration
<b>RF + all features</b>	0.74	0.80	0.45	Good discrimination

*Table 1: Model performance on test data.*

- The **Random Forest** with all features achieved the highest accuracy (≈0.74) and best AUC.
- **Logistic Regression** also outperformed Elo, achieving ~0.67 accuracy.
- The Elo-only model was notably weaker (60%), but still useful; it produces reasonable calibration (players' win probabilities are meaningful).
- Adding Elo as a feature to the ML models gave marginal gains (e.g. LR+Elo vs. LR).

**Feature Importance:** In tree models, the most important features were: Elo rating difference, head-to-head record, ranking points difference, and recent aces or first-serve win% differences. For example, if Player A had beaten Player B in their last encounter, the model leaned toward A. Players' ATP ranking points had moderate impact, reflecting current form. Surface type also mattered: separate Elo or H2H on clay vs hard sometimes changed predictions.

**Calibration:** We plotted probability reliability curves. Elo-based probabilities were somewhat conservative (often predicting closer to 50-50 than actual outcomes) due to low variability. Logistic and RF probabilities matched observed win rates closely after probability calibration (using isotonic regression).

## 11. Future Work

Our results confirm that while Elo ratings capture player strength, combining them with contextual features yields better prediction. The Random Forest’s superior accuracy suggests nonlinear interactions (e.g. a high-ranking player with many aces is very likely to win). Feature analysis indicates **serve metrics** (aces, 1st-serve points won) and **break-point saving** rates are predictive, consistent with tennis dynamics. The Elo rating difference often served as a strong single predictor, but it was enhanced by other factors.

The optimal Elo K-factor (~25) aligns with the idea that tennis has moderate volatility: too high K over-adjusts from one match. This finding is in line with sports forecasting literature (smaller K for stable ratings). The decline of Elo accuracy at high K (Figure 1) illustrates overfitting to recent matches.

We also note a gender effect: Elo predictions tended to be more accurate in women’s matches (as reported in [13]), possibly due to more stable outcomes. Although our dataset mixes ATP (men’s) matches, future work could separate men’s and women’s data to tune K differently.

## CONCLUSION

This project developed an integrated system for tennis match outcome prediction. We demonstrated that Elo ratings, while useful, benefit from being combined with machine learning on richer features. Our system achieved ~74% accuracy using a Random Forest, outperforming Elo-only (~60%). Such predictive performance can assist coaches in planning and provide bettors or commentators with probabilistic insights. Additionally, the analysis highlighted key performance drivers (Elo diff, serve stats, recent form). The methodology and code (provided in the appendix) can be extended or deployed to new seasons easily. Overall, the project shows the utility of melding traditional rating systems with modern ML for sports analytics

## REFERENCES

1. Elo, A. (1978). *The Rating of Chessplayers, Past and Present*. Arco Publishing. [en.wikipedia.org](https://en.wikipedia.org).
2. Vaughan Williams, L. et al. (2021). "How well do Elo-based ratings predict professional tennis matches?" *J. Quant. Anal. Sports*, 17(2), 91–105. DOI:10.1515/jqas-2019-0110 [degruyter.com](https://degruyter.com).
3. Bunker, R. et al. (2023). *A Comparative Evaluation of Elo Ratings and Machine Learning Methods for Tennis*. Preprint. [researchgate.net](https://researchgate.net).
4. Somboonphokkaphan, A. et al. (2009). "Tennis winner prediction based on time-series history with neural modeling." *IMCES 2009 Conference Proceedings*, 18–20. [arxiv.org](https://arxiv.org).
5. Sipko, M. & Knottenbelt, W. (2015). "Machine learning for the prediction of professional tennis matches." Imperial College MEng thesis. [arxiv.org](https://arxiv.org).
6. Hunter, J. D. (2007). *Matplotlib: A 2D graphics environment*. Computing in Science & Engineering, 9(3):90–95. [bircu-journal.com](https://bircu-journal.com).
7. Pedregosa, F. et al. (2011). *Scikit-learn: Machine Learning in Python*. Journal of Machine Learning Research, 12, 2825–2830. [jmlr.org](https://jmlr.org).
8. McKinney, W. (2010). *Data Structures for Statistical Computing in Python*. Proc. 9th Python in Science Conf. (citation for Pandas).

